

Housing Analysis

Author: Eric Wehmüller

Overview

This project is the second project for Flatiron School's bootcamp program in Data Science. We are being placed into a hypothetical situation as a Data Scientist and hoping to provide value to our business for the scenario we are given.

Business Problem

I have been hired by a real estate agency that helps homeowners sell homes. For this project, I am to provide expected/estimated home prices to homeowners based on the logistics of their home. This can also give insight on how home renovations might increase the estimated value of their homes, and what type of potential renovations are best.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
import scipy.stats as stats
%matplotlib inline

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from statsmodels.formula.api import ols
from sklearn.preprocessing import StandardScaler
from statsmodels.stats.outliers_influence import varianc
```

Data Investigation and Cleaning

To start, we have access to the King County House Sales dataset. Let's take a look at this to get a feel for what our starting point is and what raw data we have to work with.

```
df_original = pd.read_csv("data\kc_house_data.csv")
```

```
df_original.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column              Non-Null Count  Dtype  
---  -
0   id                   21597 non-null  int64  
1   date                 21597 non-null  object  
2   price                21597 non-null  float64 
3   bedrooms             21597 non-null  int64  
4   bathrooms            21597 non-null  float64 
5   sqft_living          21597 non-null  int64  
6   sqft_lot             21597 non-null  int64  
7   floors               21597 non-null  float64 
8   waterfront           19221 non-null  float64 
9   view                 21534 non-null  float64 
10  condition            21597 non-null  int64  
11  grade                21597 non-null  int64  
12  sqft_above           21597 non-null  int64  
13  sqft_basement        21597 non-null  object  
14  yr_built              21597 non-null  int64  
15  yr_renovated         17755 non-null  float64 
16  zipcode              21597 non-null  int64  
17  lat                  21597 non-null  float64 
18  long                 21597 non-null  float64 
19  sqft_living15        21597 non-null  int64  
20  sqft_lot15           21597 non-null  int64  
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB
```

```
df_original.head(10)
```

	id	date	price	bedrooms	bathrooms
0	7129300520	10/13/2014	221900.0	3	1.00
1	6414100192	12/9/2014	538000.0	3	2.25
2	5631500400	2/25/2015	180000.0	2	1.00
3	2487200875	12/9/2014	604000.0	4	3.00
4	1954400510	2/18/2015	510000.0	3	2.00
5	7237550310	5/12/2014	1230000.0	4	4.50
6	1321400060	6/27/2014	257500.0	3	2.25
7	2008000270	1/15/2015	291850.0	3	1.50
8	2414600126	4/15/2015	229500.0	3	1.00
9	3793500160	3/12/2015	323000.0	3	2.50

10 rows x 21 columns

Per the project description, I will be ignoring the following features: date, view, sqft_above, sqft_basement, yr_renovated, zipcode, lat, long, sqft_living15, sqft_lot15. For the time being, I am trying to make my modeling phase in this project as simple as possible.

```
df_col_drops = df_original.drop(columns=['id', 'date', 'view', 'yr_renovated', 'zipcode', 'lat', 'long', 'sqft_living15', 'sqft_lot15'])
display(df_col_drops)
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot
0	221900.0	3	1.00	1180	565
1	538000.0	3	2.25	2570	724
2	180000.0	2	1.00	770	100
3	604000.0	4	3.00	1960	500
4	510000.0	3	2.00	1680	808
...
21592	360000.0	3	2.50	1530	113
21593	400000.0	4	2.50	2310	581
21594	402101.0	2	0.75	1020	135
21595	400000.0	3	2.50	1600	238
21596	325000.0	2	0.75	1020	107

21597 rows x 10 columns

```
df_col_drops.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   price            21597 non-null  float64
1   bedrooms         21597 non-null  int64
2   bathrooms        21597 non-null  float64
3   sqft_living      21597 non-null  int64
4   sqft_lot         21597 non-null  int64
5   floors           21597 non-null  float64
6   waterfront       19221 non-null  float64
7   condition        21597 non-null  int64
8   grade            21597 non-null  int64
9   yr_built         21597 non-null  int64
dtypes: float64(4), int64(6)
memory usage: 1.6 MB
```

Waterfront appears to have ~2000 null values. Let's investigate what values are in this column to see what we can do about the null values.

Which ones are the most important features?

```
df_col_drops.waterfront.value_counts()
```

```
0.0    19075
1.0      146
Name: waterfront, dtype: int64
```

Only 146 have a waterfront view. Since this is a binary-filled column, I believe we can fill in all NaNs with a zero value. This makes sense, as NaNs almost certainly denotes the absence of a waterfront view.

```
df_col_drops.waterfront.fillna(0, inplace=True)
display(df_col_drops.head())
```

	price	bedrooms	bathrooms	sqft_living	sqft_lo
0	221900.0	3	1.00	1180	5650
1	538000.0	3	2.25	2570	7242
2	180000.0	2	1.00	770	10000
3	604000.0	4	3.00	1960	5000
4	510000.0	3	2.00	1680	8080

```
df_col_drops.describe()
```


	price	bedrooms	bathrooms	sqft_liv
count	2.159700e+04	21597.000000	21597.000000	21597.000000
mean	5.402966e+05	3.373200	2.115826	2080.3218
std	3.673681e+05	0.926299	0.768984	918.10612
min	7.800000e+04	1.000000	0.500000	370.00000
25%	3.220000e+05	3.000000	1.750000	1430.0000
50%	4.500000e+05	3.000000	2.250000	1910.0000
75%	6.450000e+05	4.000000	2.500000	2550.0000
max	7.700000e+06	33.000000	8.000000	13540.0000

```
df_col_drops.columns
```

```
Index(['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'waterfront', 'condition', 'grade', 'yr_built'], dtype='object')
```

```
#iterating over all columns except id to see general dis
```

```
df_col_drops.hist(figsize = (20,15));
```



It appears that we have some outliers in this data, so it's a little difficult to get a sense for what some the distributions actually are.

Specifically, I'm seeing a single entry priced at 7.7 million.



I also can't really tell what the bedroom distribution is with an outlier of 33.



sqft_lot has only a single column in this view and the mean is vastly different from the median. We will need to take a closer look at this as well.

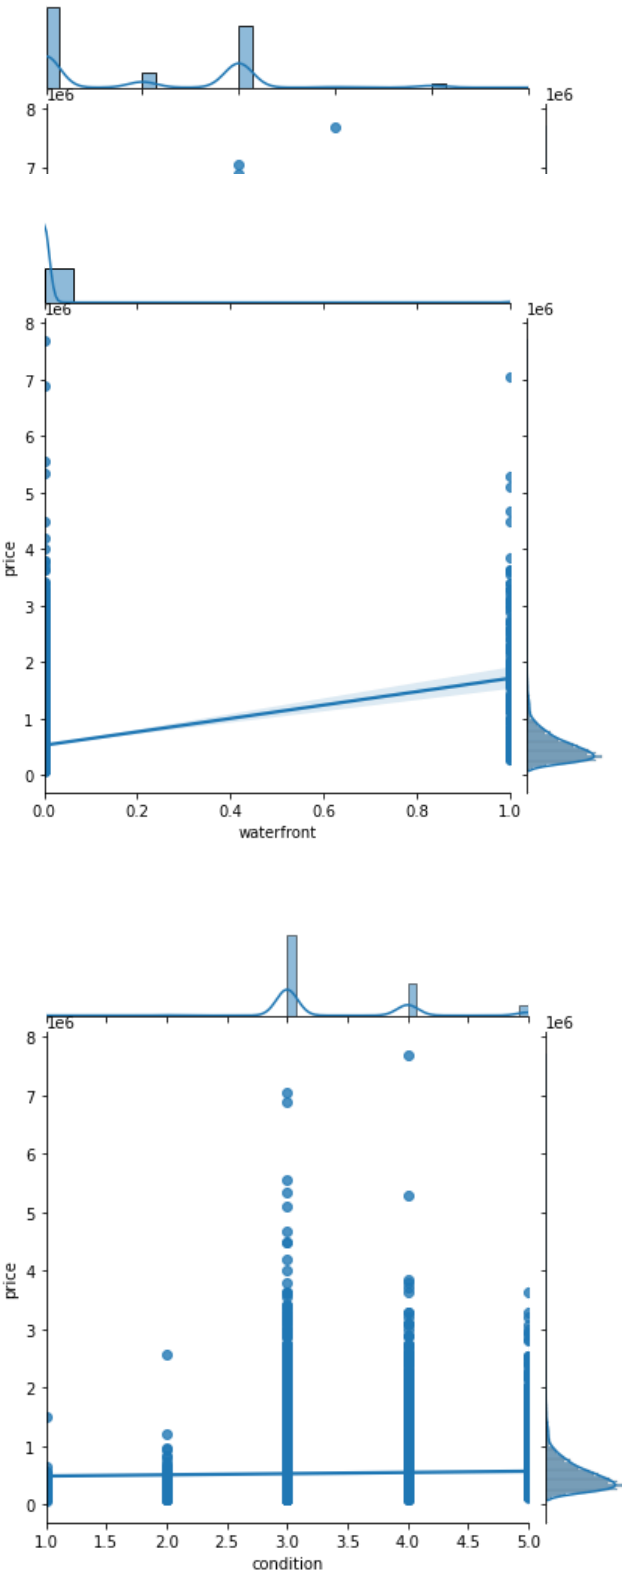
Condition and grade seem to be relatively normal.



```
#Check for linearity via jointplots
```

```
for col_name in df_col_drops.columns[1:]:  
    sns.jointplot(x=col_name, y='price', data=df_col_dro
```









It worth noting that these jointplots reveal several of these columns to have linear relations with price.

Strong Linear Relation: sqft_living, grade

Somehwat Linear: bathrooms, sqft_lot, waterfront

Little to No Linear Relation: bedrooms, floors, condition, yr_built

It appears that the features that have the largest impact on the price of a home are the square footage of the home, as well as the Grade- this rating is given by the King County Housing System. I have copied this system below for more context.

1-3 Falls short of minimum building standards. Normally cabin or inferior structure.

4 Generally older, low quality construction. Does not meet code.

5 Low construction costs and workmanship. Small, simple design.

6 Lowest grade currently meeting building code. Low quality materials and simple designs.

7 Average grade of construction and design. Commonly seen in plats and older sub-divisions.

8 Just above average in construction and design. Usually better materials in both the exterior and interior finish work.

9 Better architectural design with extra interior and exterior design and quality.

10 Homes of this quality generally have high quality features. Finish work is better and more design quality is seen in the floor plans. Generally have a larger square footage.

11 Custom design and higher quality finish work with added amenities of solid woods, bathroom fixtures and more luxurious options.

12 Custom design and excellent builders. All materials are of the highest quality and all conveniences are present.

13 Generally custom designed and built. Mansion level. Large amount of highest quality cabinet work, wood trim, marble, entry ways etc.

Modeling

Model 1

```
outcome = 'price'
x_cols = list(df_col_drops.columns)
x_cols.remove(outcome)
print(x_cols)
```

```
['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'waterfront', 'condition', 'grade', 'yr_built']
```

```
train, test = train_test_split(df_col_drops)
```

```
for col in x_cols:
    train[col] = (train[col] - train[col].mean())/train[
display(train.head())
print(len(train), len(test))
```

<ipython-input-15-f07e438ec62e>:2: SettingWithCopyWarning:
g:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
train[col] = (train[col] - train[col].mean())/train[col].std()
```

	price	bedrooms	bathrooms	sqft_living	sqft_above
12536	442500.0	-0.402234	-0.803642	0.682700	0.711700
13098	381000.0	-0.402234	-0.479047	0.671829	0.000000
17379	440000.0	-0.402234	-0.154452	-0.317497	0.400000
14127	307300.0	-1.465211	-0.154452	-0.611033	-0.300000
14618	267000.0	0.660743	-0.479047	-0.089191	-0.100000

```
16197 5400
```

```

predictors = '+' .join(x_cols)
formula = outcome + '~' + predictors
model = ols(formula=formula, data=train).fit()
model.summary()

```

OLS Regression Results


Dep. Variable:	price	R-squared:	0.646		
Model:	OLS	Adj. R-squared:	0.646		
Method:	Least Squares	F-statistic:	3284.		
Date:	Fri, 19 Mar 2021	Prob (F-statistic):	0.00		
Time:	10:05:27	Log-Likelihood:	-2.2218e+05		
No. Observations:	16197	AIC:	4.444e+05		
Df Residuals:	16187	BIC:	4.445e+05		
Df Model:	9				
Covariance Type:	nonrobust				
	coef	std err	t	P> t	[0.025
Intercept	5.403e+05	1724.545	313.278	0.000	5.37e+05
bedrooms	-3.696e+04	2185.606	-16.911	0.000	-4.12e+04
bathrooms	3.902e+04	3084.654	12.650	0.000	3.3e+04
sqft_living	1.609e+05	3510.825	45.828	0.000	1.54e+05
sqft_lot	-9372.8854	1766.156	-5.307	0.000	-1.28e+04
floors	8600.5384	2161.040	3.980	0.000	4364.661
waterfront	6.307e+04	1742.891	36.184	0.000	5.96e+04
condition	1.199e+04	1876.567	6.387	0.000	8307.926
grade	1.542e+05	2918.802	52.833	0.000	1.48e+05
yr_built	-1.117e+05	2266.720	-49.293	0.000	-1.16e+05
Omnibus:	12304.793	Durbin-Watson:	2.016		
Prob(Omnibus):	0.000	Jarque-Bera (JB):	895270.137		
Skew:	3.055	Prob(JB):	0.00		
Kurtosis:	38.906	Cond. No.	4.74		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


The p-values are less than 0.05 for our selected columns.
Let's take a look at our residuals for normality.

```
fig = sm.graphics.qqplot(model.resid, dist=stats.norm, l
```



```
plt.scatter(model.predict(train[x_cols]), model.resid)
plt.plot(model.predict(train[x_cols]), [0 for i in range
```

```
[<matplotlib.lines.Line2D at 0x1f955244820>]
```



This doesn't look great, as our QQ plot looks incorrect and we have a pronounced funnel shape on our check for homoscedasticity. We are going to need to make some changes.


Model 2 Iterations

For this iteration, I'm going to remove some outliers. (log transformation?)

I recall having the most issues determining the normal distributions of sqft_lot and bedrooms, so I'm going to filter on both.

```
df_col_drops.sqft_lot.hist()
```

<AxesSubplot:>




```
for i in range(80,100):
    q = i/100
    print("{} percentile: {}".format(q,df_col_drops.sqft_
```

```
0.8 percentile: 12182.399999999998
0.81 percentile: 12558.0
0.82 percentile: 13055.439999999995
0.83 percentile: 13503.68
0.84 percentile: 14197.0
0.85 percentile: 15000.0
0.86 percentile: 15716.040000000012
0.87 percentile: 16646.640000000003
0.88 percentile: 18000.0
0.89 percentile: 19550.0
0.9 percentile: 21371.600000000006
0.91 percentile: 24149.360000000015
0.92 percentile: 28505.119999999995
0.93 percentile: 34848.0
0.94 percentile: 37643.19999999999
0.95 percentile: 43307.200000000026
0.96 percentile: 50655.28
0.97 percentile: 67381.71999999999
0.98 percentile: 107157.0
0.99 percentile: 213008.0
```

I think filtering out homes with greater than 100k square feet is acceptable here.

```
df_col_drops.bedrooms.hist()
```

<AxesSubplot:>



```
for i in range(80,100):  
    q = i/100  
    print("{} percentile: {}".format(q,df_col_drops.bedr
```

```
0.8 percentile: 4.0  
0.81 percentile: 4.0  
0.82 percentile: 4.0  
0.83 percentile: 4.0  
0.84 percentile: 4.0  
0.85 percentile: 4.0  
0.86 percentile: 4.0  
0.87 percentile: 4.0  
0.88 percentile: 4.0  
0.89 percentile: 4.0  
0.9 percentile: 4.0  
0.91 percentile: 4.0  
0.92 percentile: 5.0  
0.93 percentile: 5.0  
0.94 percentile: 5.0  
0.95 percentile: 5.0  
0.96 percentile: 5.0  
0.97 percentile: 5.0  
0.98 percentile: 5.0  
0.99 percentile: 6.0
```



```
df_col_drops.bedrooms.value_counts()
```

3	9824
4	6882
2	2760
5	1601
6	272
1	196
7	38
8	13
9	6
10	3
11	1
33	1

Name: bedrooms, dtype: int64

I will also be filtering out all houses with more than 6 bedrooms, removing about 2% of the total entries. (may overlap with sq footage)

I will also include a log transformation to the price feature, as this may help fix our QQplot from Model 1.

```

orig_tot = len(df_col_drops)
df_outlier_filter = df_col_drops.copy()
df_outlier_filters = df_outlier_filter[df_outlier_filter
print('Percent removed:', (orig_tot -len(df_outlier_filt

df_outlier_filters = df_outlier_filters[df_outlier_filt
print('Percent removed:', (orig_tot -len(df_outlier_filt

#applying a log transformation to the price, which is ri
df_outlier_filter['price'] = np.log(df_outlier_filter['p

train2, test2 = train_test_split(df_outlier_filters)

# Refit model with subset features
predictors = '+'.join(x_cols)
formula = outcome + "~" + predictors
model2 = ols(formula=formula, data=train2).fit()
model2.summary()

```

Percent removed: 0.021530768162244755

Percent removed: 0.024355234523313424

OLS Regression Results

Dep. Variable:	price	R-squared:	0.652		
Model:	OLS	Adj. R-squared:	0.652		
Method:	Least Squares	F-statistic:	3284.		
Date:	Fri, 19 Mar 2021	Prob (F-statistic):	0.00		
Time:	10:05:28	Log-Likelihood:	-2.1609e+05		
No. Observations:	15803	AIC:	4.322e+05		
Df Residuals:	15793	BIC:	4.323e+05		
Df Model:	9				
Covariance Type:	nonrobust				
	coef	std err	t	P> t	[0.025
Intercept	6.539e+06	1.46e+05	44.654	0.000	6.25e+06
bedrooms	-4.765e+04	2485.774	-19.169	0.000	-5.25e+04
bathrooms	5.407e+04	3930.891	13.756	0.000	4.64e+04
sqft_living	179.6059	3.906	45.984	0.000	171.950
sqft_lot	-1.5529	0.160	-9.714	0.000	-1.866


floors	1.155e+04	3939.660	2.932	0.003	3828.551
waterfront	7.974e+05	2.18e+04	36.615	0.000	7.55e+05
condition	2.005e+04	2784.624	7.201	0.000	1.46e+04
grade	1.304e+05	2446.460	53.311	0.000	1.26e+05
yr_built	-3753.7342	75.323	-49.835	0.000	-3901.375
Omnibus:	10197.238	Durbin-Watson:	1.997		
Prob(Omnibus):	0.000	Jarque-Bera (JB):	438980.484		
Skew:	2.514	Prob(JB):	0.00		
Kurtosis:	28.326	Cond. No.	1.34e+06		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


[2] The condition number is large, 1.34e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
fig = sm.graphics.qqplot(model2.resid, dist=stats.norm,
```



```
plt.scatter(model2.predict(train2[x_cols]), model2.resid
plt.plot(model2.predict(train2[x_cols]), [0 for i in ran
```

```
[<matplotlib.lines.Line2D at 0x1f95673ef40>]
```



Similar problems as last time, but our OLS has alerted us that there is strong collinearity. Let's investigate what we should remove.

```
X = df_col_drops[x_cols]
vif = [variance_inflation_factor(X.values, i) for i in r
list(zip(x_cols, vif))]
```

```
[('bedrooms', 23.09608478897893),
 ('bathrooms', 24.591759875968087),
 ('sqft_living', 25.181513925621946),
 ('sqft_lot', 1.18527557276325),
 ('floors', 13.133195105016583),
 ('waterfront', 1.0252421775002192),
 ('condition', 29.533165474917077),
 ('grade', 124.69739326481557),
 ('yr_built', 124.82668596464562)]
```

Going to drop 'grade' and 'yr_built' from our model for the time being and go from there. You usually want to remove

variables with a cif of 10 or greater, indicating that they are displaying multicollinearity with other variables in the feature set.

```

train2a, test2a = train_test_split(df_outlier_filter)

outcome = 'price'
x_cols = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot']
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model2a = ols(formula=formula, data=train2a).fit()
model2a.summary()

```

OLS Regression Results

Dep. Variable:	price	R-squared:	0.512			
Model:	OLS	Adj. R-squared:	0.512			
Method:	Least Squares	F-statistic:	2427.			
Date:	Fri, 19 Mar 2021	Prob (F-statistic):	0.00			
Time:	10:05:28	Log-Likelihood:	-6739.1			
No. Observations:	16197	AIC:	1.349e+04			
Df Residuals:	16189	BIC:	1.356e+04			
Df Model:	7					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	11.9256	0.022	552.526	0.000	11.883	11.968
bedrooms	-0.0612	0.004	-14.980	0.000	-0.069	-0.053
bathrooms	0.0371	0.006	5.888	0.000	0.025	0.050
sqft_living	0.0004	5.25e-06	74.830	0.000	0.000	0.000
sqft_lot	-1.69e-07	6.85e-08	-2.468	0.014	-3.03e-07	-3.48e-08
floors	0.0893	0.006	13.983	0.000	0.077	0.102
waterfront	0.5934	0.035	17.050	0.000	0.525	0.662
condition	0.0875	0.005	18.884	0.000	0.078	0.097
Omnibus:	5.730	Durbin-Watson:	1.980			
Prob(Omnibus):	0.057	Jarque-Bera (JB):	6.076			
Skew:	0.012	Prob(JB):	0.0479			
Kurtosis:	3.092	Cond. No.	5.51e+05			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.51e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
X = df_col_drops[x_cols]
vif = [variance_inflation_factor(X.values, i) for i in range(X.shape[0])]
list(zip(x_cols, vif))

[('bedrooms', 20.155807501016387),
 ('bathrooms', 24.05414728245284),
 ('sqft_living', 16.680775340246925),
 ('sqft_lot', 1.17958229220644),
 ('floors', 10.093490615100084),
 ('waterfront', 1.0251110953470544),
 ('condition', 10.865278180945477)]
```

```

train2b, test2b = train_test_split(df_outlier_filter)

x_cols = ['sqft_living', 'sqft_lot', 'floors', 'waterfront']
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model2b = ols(formula=formula, data=train2b).fit()
model2b.summary()

```

OLS Regression Results

Dep. Variable:	price	R-squared:	0.505			
Model:	OLS	Adj. R-squared:	0.505			
Method:	Least Squares	F-statistic:	3307.			
Date:	Fri, 19 Mar 2021	Prob (F-statistic):	0.00			
Time:	10:05:28	Log-Likelihood:	-6837.0			
No. Observations:	16197	AIC:	1.369e+04			
Df Residuals:	16191	BIC:	1.373e+04			
Df Model:	5					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	11.8351	0.020	588.019	0.000	11.796	11.875
sqft_living	0.0004	3.45e-06	108.001	0.000	0.000	0.000
sqft_lot	-1.919e-07	7.12e-08	-2.696	0.007	-3.31e-07	-5.24e-08
floors	0.1051	0.006	17.704	0.000	0.093	0.117
waterfront	0.6124	0.035	17.475	0.000	0.544	0.681
condition	0.0816	0.005	17.591	0.000	0.073	0.091
Omnibus:	1.639	Durbin-Watson:	2.028			
Prob(Omnibus):	0.441	Jarque-Bera (JB):	1.639			
Skew:	0.009	Prob(JB):	0.441			
Kurtosis:	2.954	Cond. No.	5.34e+05			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.34e+05. This might

indicate that there are
strong multicollinearity or other numerical problems.

```
X = df_col_drops[x_cols]
vif = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
list(zip(x_cols, vif))

[('sqft_living', 7.1722391606368525),
 ('sqft_lot', 1.1732206092964013),
 ('floors', 7.77186134275548),
 ('waterfront', 1.017459668782625),
 ('condition', 6.7043028937497136)]
```

```

train2c, test2c = train_test_split(df_outlier_filter)

x_cols = ['sqft_living', 'sqft_lot', 'waterfront', 'condition']
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model2c = ols(formula=formula, data=train2c).fit()
model2c.summary()

```

OLS Regression Results

Dep. Variable:	price	R-squared:	0.502			
Model:	OLS	Adj. R-squared:	0.502			
Method:	Least Squares	F-statistic:	4081.			
Date:	Fri, 19 Mar 2021	Prob (F-statistic):	0.00			
Time:	10:05:29	Log-Likelihood:	-7009.4			
No. Observations:	16197	AIC:	1.403e+04			
Df Residuals:	16192	BIC:	1.407e+04			
Df Model:	4					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	12.0085	0.017	689.970	0.000	11.974	12.043
sqft_living	0.0004	3.26e-06	122.372	0.000	0.000	0.000
sqft_lot	-2.595e-07	7.31e-08	-3.548	0.000	-4.03e-07	-1.16e-07
waterfront	0.6200	0.034	18.106	0.000	0.553	0.687
condition	0.0610	0.005	13.489	0.000	0.052	0.070
Omnibus:	46.777	Durbin-Watson:	1.997			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	36.592			
Skew:	0.002	Prob(JB):	1.13e-08			
Kurtosis:	2.767	Cond. No.	5.07e+05			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


[2] The condition number is large, 5.07e+05. This might

indicate that there are
strong multicollinearity or other numerical problems.

```
X = df_col_drops[x_cols]
vif = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
list(zip(x_cols, vif))
```


```
[('sqft_living', 5.2175232204298165),
 ('sqft_lot', 1.1680613101385844),
 ('waterfront', 1.0162448940537252),
 ('condition', 4.999908921746516)]
```

```
fig = sm.graphics.qqplot(model2c.resid, dist=stats.norm,
```



```
plt.scatter(model2c.predict(train2c[x_cols]), model2c.resid)
plt.plot(model2c.predict(train2c[x_cols]), [0] * len(model2c.predict(train2c[x_cols])))
```

```
[<matplotlib.lines.Line2D at 0x1f95685af70>]
```



..... This is a modeling choice. There are pros and cons to this approach versus the first model. Removing multiple components has substantially diminished the model's

performance, as indicated by the r-squared value. However, multicollinearity between the features has been reduced.

Model 3


Going back to the drawing board, let's look at a multicollinearity heatmap to determine the columns to remove from our model.

```
first_features = ['sqft_living', 'grade', 'bathrooms', 'sqft_lot']
corr = df_col_drops[first_features].corr()
corr
```

	sqft_living	grade	bathrooms	sqft_lot
sqft_living	1.000000	0.762779	0.755758	0.173453
grade	0.762779	1.000000	0.665838	0.114731
bathrooms	0.755758	0.665838	1.000000	0.088373
sqft_lot	0.173453	0.114731	0.088373	1.000000
waterfront	0.104637	0.082818	0.063629	0.021459
floors	0.353953	0.458794	0.502582	-0.004814
condition	-0.059445	-0.146896	-0.126479	-0.008830
yr_built	0.318152	0.447865	0.507173	0.052946

```
sns.heatmap(corr, center=0, annot=True)
```

<AxesSubplot:>



sqft_living and grade = 0.76

sqft_living and bathrooms = 0.76

grade and bathrooms = 0.67

Let's remove grade and bathrooms for this model. We will also use our previous outlier filter, as this seems to be a step in the right direction.

```

train3, test3 = train_test_split(df_outlier_filter)

x_cols = ['sqft_living', 'sqft_lot', 'waterfront', 'floors']
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model3 = ols(formula=formula, data=train3).fit()
model3.summary()

```

OLS Regression Results

Dep. Variable:	price	R-squared:	0.537			
Model:	OLS	Adj. R-squared:	0.537			
Method:	Least Squares	F-statistic:	3126.			
Date:	Fri, 19 Mar 2021	Prob (F-statistic):	0.00			
Time:	10:05:30	Log-Likelihood:	-6355.8			
No. Observations:	16197	AIC:	1.273e+04			
Df Residuals:	16190	BIC:	1.278e+04			
Df Model:	6					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	19.1667	0.231	83.101	0.000	18.715	19.619
sqft_living	0.0004	3.42e-06	115.938	0.000	0.000	0.000
sqft_lot	-1.651e-07	7.15e-08	-2.310	0.021	-3.05e-07	-2.5e-08
waterfront	0.5972	0.035	16.833	0.000	0.528	0.667
floors	0.1724	0.006	27.550	0.000	0.160	0.185
condition	0.0380	0.005	8.137	0.000	0.029	0.047
yr_built	-0.0037	0.000	-31.884	0.000	-0.004	-0.003
Omnibus:	115.168	Durbin-Watson:	2.011			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	141.348			
Skew:	-0.133	Prob(JB):	2.03e-31			
Kurtosis:	3.372	Cond. No.	3.52e+06			


Notes:

[1] Standard Errors assume that the covariance matrix of

the errors is correctly specified.


[2] The condition number is large, $3.52e+06$. This might indicate that there are strong multicollinearity or other numerical problems.

```
fig = sm.graphics.qqplot(model3.resid, dist=stats.norm,
```



```
plt.scatter(model3.predict(train3[x_cols]), model3.resid  
plt.plot(model3.predict(train3[x_cols]), [0 for i in ran
```

```
[<matplotlib.lines.Line2D at 0x1f94f8e4790>]
```



Model 4

Our QQ plots are less than ideal in previous models. Let's see if we can fix that by using a transform on the appropriate features.

```
for col_name in df_outlier_filter.columns[1:]:  
    print(col_name)  
    print(df_outlier_filter[col_name].skew())
```

```
bedrooms  
2.023641235344595  
bathrooms  
0.5197092816403838  
sqft_living  
1.473215455425834  
sqft_lot  
13.072603567136046  
floors  
0.6144969756263127  
waterfront  
12.039584643829357  
condition  
1.0360374245132955  
grade  
0.7882366363846076  
yr_built  
-0.4694499764949978
```

'sqft_lot' seems to be the main issue with the highest skew coefficient. I'm not sure if I should apply this to waterfront. We may need to use another method here, or look elsewhere for model improvements.


```
#only run once  
df_outlier_filter['sqft_lot'] = np.log(df_outlier_filter  
df_outlier_filter['sqft_lot'].skew())
```

```
0.9625003856495555
```



```
df_outlier_filter['sqft_lot'].hist()
```

<AxesSubplot:>



```
df_outlier_filter['bedrooms'] = np.log(df_outlier_filter  
df_outlier_filter['bedrooms'].skew())
```

-0.6805637280656164

```

x_cols = list(df_outlier_filter.columns)
x_cols.remove(outcome)

train4, test4 = train_test_split(df_outlier_filter)

predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model4 = ols(formula=formula, data=train4).fit()
model4.summary()

```

OLS Regression Results


Dep. Variable:	price	R-squared:	0.646			
Model:	OLS	Adj. R-squared:	0.645			
Method:	Least Squares	F-statistic:	3275.			
Date:	Fri, 19 Mar 2021	Prob (F-statistic):	0.00			
Time:	10:05:31	Log-Likelihood:	-4151.9			
No. Observations:	16197	AIC:	8324.			
Df Residuals:	16187	BIC:	8401.			
Df Model:	9					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	22.0160	0.214	102.848	0.000	21.596	22.436
bedrooms	-0.1065	0.011	-9.635	0.000	-0.128	-0.085
bathrooms	0.0808	0.006	14.091	0.000	0.070	0.092
sqft_living	0.0002	5.7e-06	36.097	0.000	0.000	0.000
sqft_lot	-0.0470	0.003	-14.600	0.000	-0.053	-0.041
floors	0.0520	0.006	8.578	0.000	0.040	0.064
waterfront	0.4809	0.031	15.731	0.000	0.421	0.541
condition	0.0458	0.004	11.110	0.000	0.038	0.054
grade	0.2273	0.004	63.858	0.000	0.220	0.234
yr_built	-0.0056	0.000	-50.456	0.000	-0.006	-0.005
Omnibus:	66.219	Durbin-Watson:	2.038			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	77.043			
Skew:	-0.100	Prob(JB):	1.86e-17			
Kurtosis:	3.273	Cond. No.	2.57e+05			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


[2] The condition number is large, $2.57\text{e}+05$. This might indicate that there are strong multicollinearity or other numerical problems.

```
fig = sm.graphics.qqplot(model4.resid, dist=stats.norm,
```



```
plt.scatter(model4.predict(train4[x_cols]), model4.resid
plt.plot(model4.predict(train4[x_cols]), [0 for i in ran
```

```
[<matplotlib.lines.Line2D at 0x1f95378cdc0>]
```



This is a nice improvement. This is our best model thus far. It passes the normality check from looking at the QQ plot and it is homoscedastic.

Interpreting this model:


R-squared: 64.6% variation in the price can be explained by 'sqft_living', 'grade', 'bathrooms', 'sqft_lot', 'waterfront', 'floors', 'condition', 'yr_built', and 'bedrooms'

Durbin-watson: A value preferred between 1-2 implies that the regression results are reliable from the side of homoscedasticity.

The highest coefficients belong to Grade and Waterfront: namely, what grade the home has been given by the King County Housing System. Additionally, having a waterfront view as a part of your home largely impacts the price.

When needed, we can now use this model to give us prediction values for an estimated price, given the values for the features of a home we are trying to sell. Obviously, someone would be unable to renovate their home to suddenly have a waterfront view, but doing something like adding a bathroom (the 3rd highest coefficient) seems to also have a significant impact of the expected price of a home for this model as well.

```
model4.params[1:].abs().plot.bar(color=['red', 'blue', 'bl
```



This is a visualization of our coefficients. To compare, I have taken the absolute value of each in the series, but made sure to indicate negative coefficients in red columns.

Conclusion

I believe the best model is Model 4, where the outliers have been filtered out and none of the features are removed .

Although this suffers from multicollinearity, it has an r-squared value of ~0.647, which is the most accurate model in our analysis.

I believe this is acceptable within the context of this scenario. It affects the coefficients and p-values, but it does not influence the predictions, precision of the predictions, and the statistics determining goodness of fit. Our primary goal is to have a model to make predictions for us.

To further improve this, I would use more of the columns included in the original dataset to try to increase my r-squared value and hopefully fix the QQplot issues I was having for all of my models.