

Housing Analysis

Author: Eric Wehmüller

Overview

This project is the second project for Flatiron School's bootcamp program in Data Science. We are being placed into a hypothetical situation as a Data Scientist and hoping to provide value to our business for the scenario we are given.

Business Problem

I have been hired by a real estate agency that helps homeowners sell homes. For this project, I am to provide expected/estimated home prices to homeowners based on the logistics of their home. This can also give insight on how home renovations might increase the estimated value of their homes, and what type of potential renovations are best.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
import scipy.stats as stats
%matplotlib inline

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from statsmodels.formula.api import ols
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

Data Investigation and Cleaning

To start, we have access to the King County House Sales dataset. Let's take a look at this to get a feel for what our starting point is and what raw data we have to work with.

```
df_original = pd.read_csv("data\kc_house_data.csv")
```

```
df_original.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   id                   21597 non-null  int64
1   date                 21597 non-null  object
2   price                21597 non-null  float64
3   bedrooms             21597 non-null  int64
4   bathrooms            21597 non-null  float64
5   sqft_living          21597 non-null  int64
6   sqft_lot             21597 non-null  int64
7   floors               21597 non-null  float64
8   waterfront           19221 non-null  float64
9   view                 21534 non-null  float64
10  condition            21597 non-null  int64
11  grade                21597 non-null  int64
12  sqft_above           21597 non-null  int64
13  sqft_basement        21597 non-null  object
14  yr_built              21597 non-null  int64
15  yr_renovated          17755 non-null  float64
16  zipcode               21597 non-null  int64
17  lat                   21597 non-null  float64
18  long                  21597 non-null  float64
19  sqft_living15         21597 non-null  int64
20  sqft_lot15            21597 non-null  int64
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB
```

```
df_original.head(10)
```

	id	date	price	bedrooms	bathrooms
0	7129300520	10/13/2014	221900.0	3	1.00
1	6414100192	12/9/2014	538000.0	3	2.25
2	5631500400	2/25/2015	180000.0	2	1.00
3	2487200875	12/9/2014	604000.0	4	3.00
4	1954400510	2/18/2015	510000.0	3	2.00
5	7237550310	5/12/2014	1230000.0	4	4.50
6	1321400060	6/27/2014	257500.0	3	2.25
7	2008000270	1/15/2015	291850.0	3	1.50
8	2414600126	4/15/2015	229500.0	3	1.00
9	3793500160	3/12/2015	323000.0	3	2.50

10 rows x 21 columns

Per the project description, I will be ignoring the following features: date, view, sqft_above, sqft_basement, yr_renovated, zipcode, lat, long, sqft_living15, sqft_lot15. For the time being, I am trying to make my modeling phase in this project as simple as possible.

```
df_col_drops = df_original.drop(columns=['id', 'date', 'view', 'yr_renovated', 'zipcode', 'lat', 'long', 'sqft_living15', 'sqft_lot15'])
display(df_col_drops)
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot
0	221900.0	3	1.00	1180	565
1	538000.0	3	2.25	2570	724
2	180000.0	2	1.00	770	100
3	604000.0	4	3.00	1960	500
4	510000.0	3	2.00	1680	808
...
21592	360000.0	3	2.50	1530	113
21593	400000.0	4	2.50	2310	581
21594	402101.0	2	0.75	1020	135
21595	400000.0	3	2.50	1600	238
21596	325000.0	2	0.75	1020	107

21597 rows x 12 columns

```
df_col_drops.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   price           21597 non-null  float64
1   bedrooms        21597 non-null  int64
2   bathrooms        21597 non-null  float64
3   sqft_living      21597 non-null  int64
4   sqft_lot         21597 non-null  int64
5   floors           21597 non-null  float64
6   waterfront       19221 non-null  float64
7   condition        21597 non-null  int64
8   grade            21597 non-null  int64
9   yr_built         21597 non-null  int64
10  lat              21597 non-null  float64
11  long             21597 non-null  float64
dtypes: float64(6), int64(6)
memory usage: 2.0 MB
```

Waterfront appears to have ~2000 null values. Let's investigate what values are in this column to see what we can do about the null values.

Which ones are the most important features?

```
df_col_drops.waterfront.value_counts()
```

```
0.0    19075
1.0      146
Name: waterfront, dtype: int64
```

Only 146 have a waterfront view. Since this is a binary-filled column, I believe we can fill in all NaNs with a zero value. This makes sense, as NaNs almost certainly denotes the absence of a waterfront view.

```
df_col_drops.waterfront.fillna(0, inplace=True)
display(df_col_drops.head())
```

	price	bedrooms	bathrooms	sqft_living	sqft_lo
0	221900.0	3	1.00	1180	5650
1	538000.0	3	2.25	2570	7242
2	180000.0	2	1.00	770	10000
3	604000.0	4	3.00	1960	5000
4	510000.0	3	2.00	1680	8080

```
df_col_drops.describe()
```

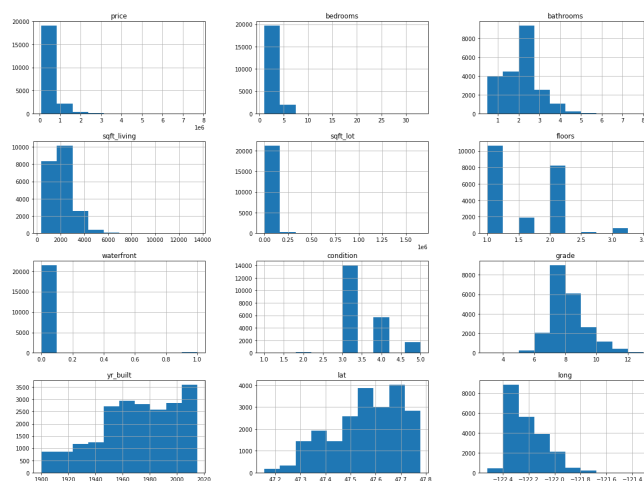
	price	bedrooms	bathrooms	sqft_liv
count	2.159700e+04	21597.000000	21597.000000	21597.000000
mean	5.402966e+05	3.373200	2.115826	2080.3218
std	3.673681e+05	0.926299	0.768984	918.10612
min	7.800000e+04	1.000000	0.500000	370.00000
25%	3.220000e+05	3.000000	1.750000	1430.0000
50%	4.500000e+05	3.000000	2.250000	1910.0000
75%	6.450000e+05	4.000000	2.500000	2550.0000
max	7.700000e+06	33.000000	8.000000	13540.0000

```
df_col_drops.columns
```

```
Index(['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',
       'waterfront', 'condition', 'grade', 'yr_built', 'lat', 'long'],
      dtype='object')
```

```
#iterating over all columns except id to see general dis
```

```
df_col_drops.hist(figsize = (20,15));
```



It appears that we have some outliers in this data, so it's a little difficult to get a sense of what some the distributions actually are.

Specifically, I'm seeing a single entry priced at 7.7 million.



I also can't really tell what the bedroom distribution is with an outlier of 33.



sqft_lot has only a single column in this view and the mean is vastly different from the median. We will need to take a closer look at this as well.

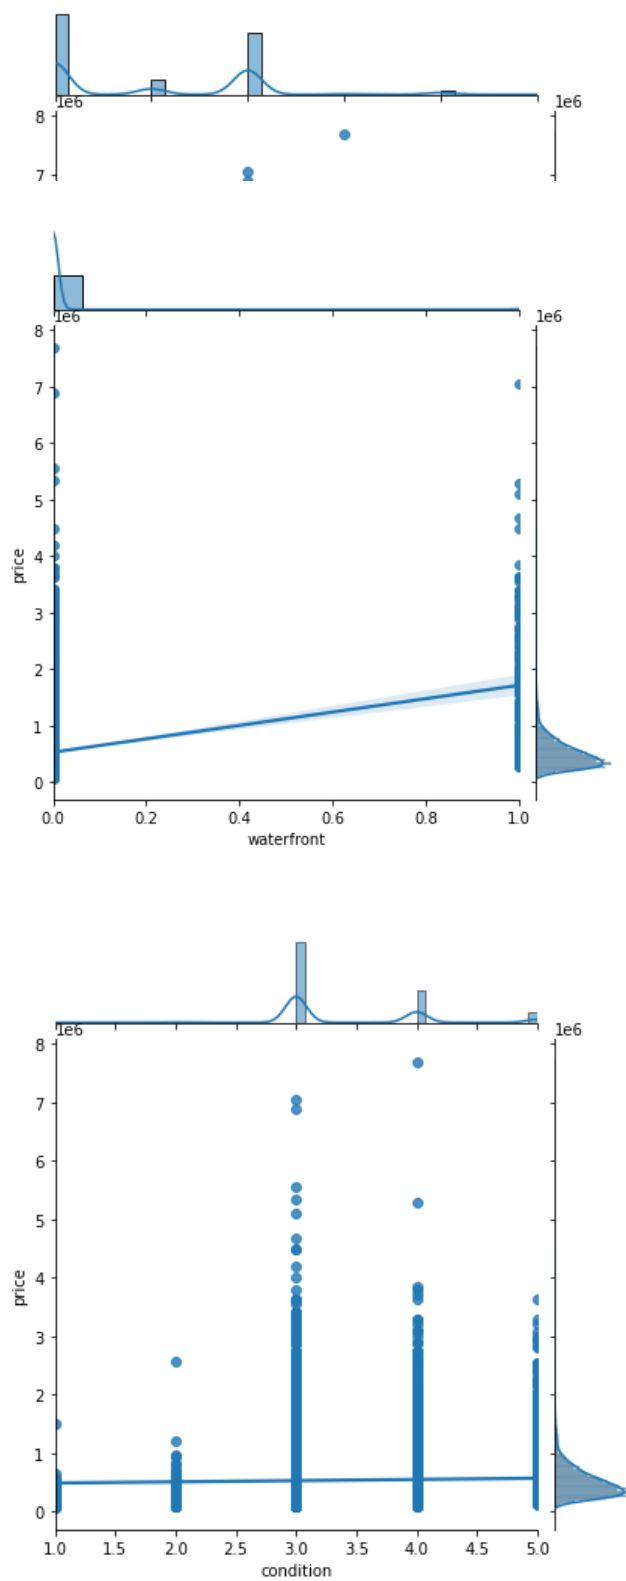
Condition and grade seem to be relatively normal.



```
#Check for linearity via jointplots
```

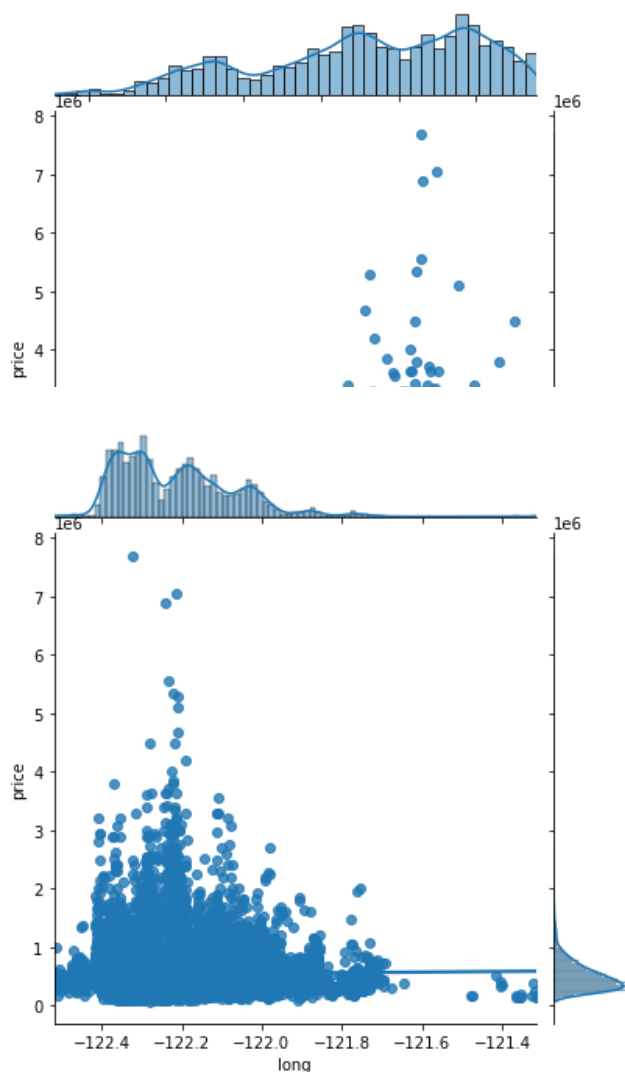
```
for col_name in df_col_drops.columns[1:]:  
    ax = sns.jointplot(x=col_name, y='price', data=df_co  
ax.
```











It worth noting that these jointplots reveal several of these columns to have linear relations with price.

Strong Linear Relation: sqft_living, grade

Somewhat Linear: bathrooms, sqft_lot, waterfront

Little to No Linear Relation: bedrooms, floors, condition, yr_built, lat, long

It appears that the features that have the largest impact on the price of a home are the square footage of the home, as well as the Grade- this rating is given by the King County Housing System. I have copied this system below for more context.

1-3 Falls short of minimum building standards. Normally cabin or inferior structure.

4 Generally older, low quality construction. Does not meet code.

5 Low construction costs and workmanship. Small, simple design.

6 Lowest grade currently meeting building code. Low quality materials and simple designs.

7 Average grade of construction and design. Commonly seen in plats and older sub-divisions.

8 Just above average in construction and design. Usually better materials in both the exterior and interior finish work.

9 Better architectural design with extra interior and exterior design and quality.

10 Homes of this quality generally have high quality features. Finish work is better and more design quality is seen in the floor plans. Generally have a larger square footage.

11 Custom design and higher quality finish work with added amenities of solid woods, bathroom fixtures and more luxurious options.

12 Custom design and excellent builders. All materials are of the highest quality and all conveniences are present.

13 Generally custom designed and built. Mansion level. Large amount of highest quality cabinet work, wood trim, marble, entry ways etc.

Feature Engineering

Two fields jump out at me: latitude and longitude. As we already know, this data is taken from the King County Housing dataset, which includes the city of Seattle. Let's engineer a feature that determines the distance from "downtown" using lat and long.

```
#using 47.605° N, 122.334° W as the exact point for downtown
dtwn_lat = 47.605
dtwn_long = -122.334
dtwn_coords = (dtwn_lat, dtwn_long)
print(type(dtwn_coords))

second_coords = (df_col_drops['lat'][0], df_col_drops['lon'][0])
print(second_coords)

<class 'tuple'>
(47.5112, -122.257)
```

```
import haversine as hs
```

```
#solving for a single location, in kilometers
hs.haversine(dtown_coords, second_coords)
```

```
11.923605090619347
```

```
#creating feature column
```

```
df_col_drops['dist_to_dtown'] = df_col_drops.lat
for index, row in df_col_drops.iterrows():
    df_col_drops['dist_to_dtown'][index] = hs.haversine(d
```

```
<ipython-input-15-edcbe8011a22>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_col_drops['dist_to_dtown'][index] = hs.haversine(dtown_coords, point2=(df_col_drops['lat'][index], df_col_drops['long'][index]))
```

```
df_col_drops.head(10)
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot
0	221900.0	3	1.00	1180	5650
1	538000.0	3	2.25	2570	7242
2	180000.0	2	1.00	770	10000
3	604000.0	4	3.00	1960	5000
4	510000.0	3	2.00	1680	8080
5	1230000.0	4	4.50	5420	101930
6	257500.0	3	2.25	1715	6819
7	291850.0	3	1.50	1060	9711
8	229500.0	3	1.00	1780	7470
9	323000.0	3	2.50	1890	6560

```
#dropping these so we don't confuse our model- dist_to_d
df_col_drops = df_col_drops.drop(['lat', 'long'], axis=1)
```

Modeling

Model 1

```
outcome = 'price'
x_cols = list(df_col_drops.columns)
x_cols.remove(outcome)
print(x_cols)

['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'waterfront', 'condition', 'grade', 'yr_built', 'dist_to_dtnw']
```

```
price_log = np.log(df_col_drops.price)
price_log = pd.DataFrame(price_log)

X1= df_col_drops.drop('price', 1)
y1= price_log
```

```
X_train, X_test, y_train, y_test = train_test_split(X1, y1,
```

```
#normalization
```

```
for col in x_cols:
    X_train[col] = (X_train[col] - X_train[col].mean())/
display(X_train.head())
print(len(X_train), len(X_test))
```

<ipython-input-21-c934d189158c>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
X_train[col] = (X_train[col] - X_train[col].mean())/X_train[col].std()
```

	bedrooms	bathrooms	sqft_living	sqft_lot	
11744	0.672806	0.508126	0.397234	-0.230020	0.9
12492	-0.397156	0.834553	-0.736429	-0.169504	-0.1
13866	-0.397156	-1.450430	-1.096150	-0.105329	-0.1
16645	-0.397156	0.508126	-0.125995	-0.155416	0.9
11548	0.672806	0.508126	-0.060591	-0.183641	0.9

```
17277 4320
```

```
# predictors = '+'.join(x_cols)
# formula = outcome + '~' + predictors
# model = ols(formula=formula, data=train).fit()
# model.summary()
```

```
predictors = sm.add_constant(X_train)
model_1 = sm.OLS(y_train, predictors).fit()
model_1.summary()
```

OLS Regression Results


Dep. Variable:	price	R-squared:	0.730			
Model:	OLS	Adj. R-squared:	0.729			
Method:	Least Squares	F-statistic:	4658.			
Date:	Mon, 29 Mar 2021	Prob (F-statistic):	0.00			
Time:	14:04:14	Log-Likelihood:	-2160.4			
No. Observations:	17277	AIC:	4343.			
Df Residuals:	17266	BIC:	4428.			
Df Model:	10					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975
const	13.0448	0.002	6251.211	0.000	13.041	13.049
bedrooms	-0.0192	0.003	-7.252	0.000	-0.024	-0.014
bathrooms	0.0434	0.004	11.617	0.000	0.036	0.050
sqft_living	0.1983	0.004	46.479	0.000	0.190	0.206
sqft_lot	0.0371	0.002	16.938	0.000	0.033	0.041
floors	0.0139	0.003	5.261	0.000	0.009	0.018
waterfront	0.0423	0.002	20.053	0.000	0.038	0.046
condition	0.0380	0.002	16.695	0.000	0.034	0.042
grade	0.2158	0.004	59.532	0.000	0.209	0.222
yr_built	-0.0564	0.003	-18.141	0.000	-0.063	-0.050
dist_to_dtnw	-0.1869	0.002	-76.113	0.000	-0.192	-0.182
Omnibus:	327.427	Durbin-Watson:	1.995			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	586.147			
Skew:	-0.144	Prob(JB):	5.25e-128			
Kurtosis:	3.855	Cond. No.	4.89			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The p-values are less than 0.05 for our selected columns. Let's take a look at our residuals for normality.

```
fig = sm.graphics.qqplot(model_1.resid, dist=stats.norm,
```



This doesn't look great, as our QQ plot looks incorrect and we have a pronounced funnel shape on our check for homoscedasticity. We are going to need to make some changes.

```
regression = LinearRegression()
regression.fit(X_train, y_train)

#use the regression for the train and test data
y_hat_train = regression.predict(X_train)
y_hat_test = regression.predict(X_test)

#Root Mean Square Error
train_rmse = np.sqrt(mean_squared_error(y_train, y_hat_train))
test_rmse = np.sqrt(mean_squared_error(y_test, y_hat_test))

print(f'Train Root Mean Square Error: {train_rmse}')
print(f'Test Root Mean Square Error: {test_rmse}')
```

```
Train Root Mean Square Error: 0.2742010910106044
Test Root Mean Square Error: 1764.4910958009116
```


Models Addressing Multicollinearity

For this iteration, I'm going to remove some outliers. (log transformation?)

I recall having the most issues determining the normal distributions of sqft_lot and bedrooms, so I'm going to filter on both.

```
df_col_drops.sqft_lot.hist()
```

<AxesSubplot:>




```
for i in range(80,100):
    q = i/100
    print("{} percentile: {}".format(q,df_col_drops.sqft_

0.8 percentile: 12182.399999999998
0.81 percentile: 12558.0
0.82 percentile: 13055.439999999995
0.83 percentile: 13503.68
0.84 percentile: 14197.0
0.85 percentile: 15000.0
0.86 percentile: 15716.040000000012
0.87 percentile: 16646.640000000003
0.88 percentile: 18000.0
0.89 percentile: 19550.0
0.9 percentile: 21371.600000000006
0.91 percentile: 24149.360000000015
0.92 percentile: 28505.119999999995
0.93 percentile: 34848.0
0.94 percentile: 37643.19999999999
0.95 percentile: 43307.200000000026
0.96 percentile: 50655.28
0.97 percentile: 67381.71999999999
0.98 percentile: 107157.0
0.99 percentile: 213008.0
```

I think filtering out homes with greater than 100k square feet is acceptable here.

```
df_col_drops.bedrooms.hist()
```

<AxesSubplot:>



```
for i in range(80,100):  
    q = i/100  
    print("{} percentile: {}".format(q,df_col_drops.bedr
```

```
0.8 percentile: 4.0  
0.81 percentile: 4.0  
0.82 percentile: 4.0  
0.83 percentile: 4.0  
0.84 percentile: 4.0  
0.85 percentile: 4.0  
0.86 percentile: 4.0  
0.87 percentile: 4.0  
0.88 percentile: 4.0  
0.89 percentile: 4.0  
0.9 percentile: 4.0  
0.91 percentile: 4.0  
0.92 percentile: 5.0  
0.93 percentile: 5.0  
0.94 percentile: 5.0  
0.95 percentile: 5.0  
0.96 percentile: 5.0  
0.97 percentile: 5.0  
0.98 percentile: 5.0  
0.99 percentile: 6.0
```

```
df_col_drops.bedrooms.value_counts()
```

```

3      9824
4      6882
2      2760
5      1601
6       272
1       196
7        38
8         13
9          6
10         3
11         1
33         1

```

```
Name: bedrooms, dtype: int64
```

I will also be filtering out all houses with more than 6 bedrooms, removing about 2% of the total entries. (may overlap with sq footage)

I will also include a log transformation to the price feature, as this may help fix our QQplot from Model 1.

```

orig_tot = len(df_col_drops)
df_outlier_filter = df_col_drops.copy()
df_outlier_filters = df_outlier_filter[df_outlier_filter
print('Percent removed:', (orig_tot - len(df_outlier_filt

df_outlier_filters = df_outlier_filters[df_outlier_filt
print('Percent removed:', (orig_tot - len(df_outlier_filt

#applying a log transformation to the price, which is ri
df_outlier_filter['price'] = np.log(df_outlier_filter['p

#train2, test2 = train_test_split(df_outlier_filters)

```

```
Percent removed: 0.021530768162244755
```

```
Percent removed: 0.024355234523313424
```

```

X2 = df_outlier_filter.drop('price', 1)
y2 = df_outlier_filter['price']
X_train2, X_test2, y_train2, y_test2 = train_test_split(

# Refit model with subset features
predictors = sm.add_constant(X_train2)
model_2 = sm.OLS(y_train2, predictors).fit()
model_2.summary()

```

OLS Regression Results


Dep. Variable:	price		R-squared:		0.731	
Model:	OLS		Adj. R-squared:		0.730	
Method:	Least Squares		F-statistic:		4682.	
Date:	Mon, 29 Mar 2021		Prob (F-statistic):		0.00	
Time:	14:04:15		Log-Likelihood:		-2144.8	
No. Observations:	17277		AIC:		4312.	
Df Residuals:	17266		BIC:		4397.	
Df Model:	10					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	15.1190	0.204	74.281	0.000	14.720	15.518
bedrooms	-0.0202	0.003	-6.768	0.000	-0.026	-0.014
bathrooms	0.0559	0.005	11.456	0.000	0.046	0.065
sqft_living	0.0002	4.66e-06	46.015	0.000	0.000	0.000
sqft_lot	9.24e-07	5.21e-08	17.726	0.000	8.22e-07	1.03e-06
floors	0.0265	0.005	5.429	0.000	0.017	0.036
waterfront	0.5214	0.026	20.186	0.000	0.471	0.572
condition	0.0593	0.004	16.905	0.000	0.052	0.066
grade	0.1835	0.003	59.469	0.000	0.177	0.190
yr_built	-0.0020	0.000	-18.661	0.000	-0.002	-0.000
dist_to_dtnw	-0.0174	0.000	-74.853	0.000	-0.018	-0.016
Omnibus:	305.749	Durbin-Watson:		1.998		
Prob(Omnibus):	0.000	Jarque-Bera (JB):		533.927		
Skew:	-0.140	Prob(JB):		1.15e-116		
Kurtosis:	3.815	Cond. No.		4.39e+06		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, $4.39\text{e}+06$. This might indicate that there are strong multicollinearity or other numerical problems.

```
fig = sm.graphics.qqplot(model_2.resid, dist=stats.norm,
```



```
regression2 = LinearRegression()
regression2.fit(X_train2, y_train2)

#use the regression for the train and test data
y_hat_train2 = regression2.predict(X_train2)
y_hat_test2 = regression2.predict(X_test2)

#Root Mean Square Error
train_rmse2 = np.sqrt(mean_squared_error(y_train2, y_hat_train2))
test_rmse2 = np.sqrt(mean_squared_error(y_test2, y_hat_test2))

print(f'Train Root Mean Square Error: {train_rmse2}')
print(f'Test Root Mean Square Error: {test_rmse2}')
```

```
Train Root Mean Square Error: 0.27395412776536393
Test Root Mean Square Error: 0.274568980060669
```

Similar problems as last time, but our OLS has alerted us that there is strong collinearity. Let's investigate what we should remove.

```
X = df_col_drops[x_cols]
X['constant'] = np.ones(X.shape[0])
vif = [variance_inflation_factor(X.values, i) for i in range(X.shape[0])]
list(zip(x_cols, vif))

[('bedrooms', 1.6311136630472653),
 ('bathrooms', 3.215688966233134),
 ('sqft_living', 4.211173170919058),
 ('sqft_lot', 1.1150457494029746),
 ('floors', 1.6055072567823685),
 ('waterfront', 1.0219826889310346),
 ('condition', 1.1874067595264461),
 ('grade', 3.0015991231227797),
 ('yr_built', 2.240108337334204),
 ('dist_to_dtn', 1.4004289073409446)]
```

You usually want to remove variables with a vif of 5~10 or greater, indicating that they are displaying multicollinearity with other variables in the feature set. None of these values are really in that range.

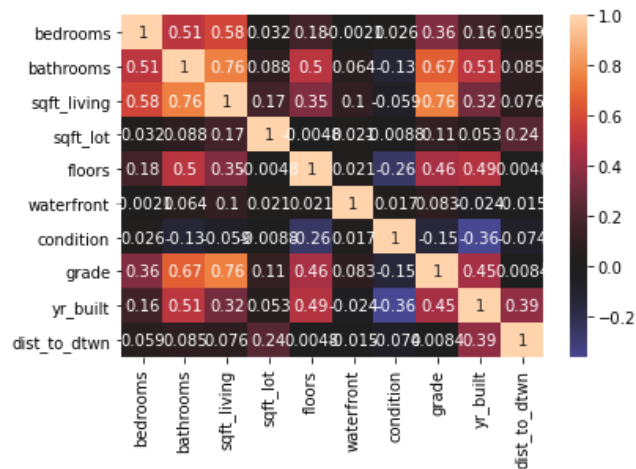
Going back to the drawing board, let's look at a multicollinearity heatmap to determine the columns to remove from our model.

```
first_features = ['bedrooms', 'bathrooms', 'sqft_living']
corr = df_col_drops[first_features].corr()
corr
```

	bedrooms	bathrooms	sqft_living	sqft_lot
bedrooms	1.000000	0.514508	0.578212	0.032471
bathrooms	0.514508	1.000000	0.755758	0.088373
sqft_living	0.578212	0.755758	1.000000	0.173453
sqft_lot	0.032471	0.088373	0.173453	1.000000
floors	0.177944	0.502582	0.353953	-0.004811
waterfront	-0.002127	0.063629	0.104637	0.021451
condition	0.026496	-0.126479	-0.059445	-0.008811
grade	0.356563	0.665838	0.762779	0.114731
yr_built	0.155670	0.507173	0.318152	0.052941
dist_to_dtn	0.058718	0.084731	0.076442	0.243471

```
sns.heatmap(corr, center=0, annot=True)
```

<AxesSubplot:>



sqft_living and grade = 0.76

sqft_living and bathrooms = 0.76

grade and bathrooms = 0.67

Let's remove grade and bathrooms for this model. We will also use our previous outlier filter, as this seems to be a step in the right direction.


```
# train3, test3 = train_test_split(df_outlier_filter)

# x_cols = ['bedrooms', 'sqft_living', 'sqft_lot', 'floors']
# predictors = '+'.join(x_cols)
# formula = outcome + '~' + predictors
# model3 = ols(formula=formula, data=train3).fit()
# model3.summary()
X3 = df_outlier_filter.drop(columns=['price', 'grade', 'bathrooms'])
y3 = df_outlier_filter['price']
X_train3, X_test3, y_train3, y_test3 = train_test_split(X3, y3,
                                                         test_size=0.2,
                                                         random_state=42)

# Refit model with subset features
predictors = sm.add_constant(X_train3)
model_3 = sm.OLS(y_train3, predictors).fit()
model_3.summary()
```

OLS Regression Results

Dep. Variable:	price	R-squared:	0.672			
Model:	OLS	Adj. R-squared:	0.672			
Method:	Least Squares	F-statistic:	4430.			
Date:	Mon, 29 Mar 2021	Prob (F-statistic):	0.00			
Time:	14:04:17	Log-Likelihood:	-3788.4			
No. Observations:	17277	AIC:	7595.			
Df Residuals:	17268	BIC:	7665.			
Df Model:	8					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	11.9934	0.208	57.584	0.000	11.585	12.401
bedrooms	-0.0413	0.003	-13.012	0.000	-0.047	-0.035
sqft_living	0.0004	3.42e-06	120.315	0.000	0.000	0.000
sqft_lot	1.108e-06	6.23e-08	17.801	0.000	9.86e-07	1.230e-06
floors	0.0855	0.005	16.478	0.000	0.075	0.095
waterfront	0.5349	0.028	18.923	0.000	0.480	0.590
condition	0.0659	0.004	17.150	0.000	0.058	0.073
yr_built	0.0002	0.000	1.656	0.098	-3.24e-05	0.000
dist to dtwn	-0.0207	0.000	-82.778	0.000	-0.021	-0.020


Omnibus:	366.331	Durbin-Watson:	2.013
Prob(Omnibus):	0.000	Jarque-Bera (JB):	496.038
Skew:	-0.261	Prob(JB):	1.94e-108
Kurtosis:	3.645	Cond. No.	3.79e+06

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.79e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
fig = sm.graphics.qqplot(model_3.resid, dist=stats.norm,
```



```
regression3 = LinearRegression()
regression3.fit(X_train3, y_train3)

#use the regression for the train and test data
y_hat_train3 = regression3.predict(X_train3)
y_hat_test3 = regression3.predict(X_test3)

#Root Mean Square Error
train_rmse3 = np.sqrt(mean_squared_error(y_train3, y_hat_train3))
test_rmse3 = np.sqrt(mean_squared_error(y_test3, y_hat_test3))

print(f'Train Root Mean Square Error: {train_rmse3}')
print(f'Test Root Mean Square Error: {test_rmse3}')
```

```
Train Root Mean Square Error: 0.3012951581957556
Test Root Mean Square Error: 0.3076316639605217
```

This is a modeling choice. There are pros and cons to this approach versus the first model. Removing multiple components has substantially diminished the model's performance, as indicated by the r-squared value. However, multicollinearity between the features has been reduced.

Model 4

Our QQ plots are less than ideal in previous models. Let's see if we can fix that by using a transform on the appropriate features.

```
for col_name in df_outlier_filter.columns[1:]:  
    print(col_name)  
    print(df_outlier_filter[col_name].skew())
```


```
bedrooms  
2.023641235344595  
bathrooms  
0.5197092816403838  
sqft_living  
1.473215455425834  
sqft_lot  
13.072603567136046  
floors  
0.6144969756263127  
waterfront  
12.039584643829357  
condition  
1.0360374245132955  
grade  
0.7882366363846076  
yr_built  
-0.4694499764949978  
dist_to_dtn  
0.769367697269784
```

'sqft_lot' seems to be the main issue with the highest skew coefficient. I'm not sure if I should apply this to waterfront. We may need to use another method here, or look elsewhere for model improvements.

```
#only run once  
df_outlier_filter['sqft_lot'] = np.log(df_outlier_filter  
df_outlier_filter['sqft_lot'].skew()  
  
0.9625003856495555
```

```
df_outlier_filter['sqft_lot'].hist()
```

<AxesSubplot:>



```
df_outlier_filter['bedrooms'] = np.log(df_outlier_filter  
df_outlier_filter['bedrooms'].skew())
```

-0.6805637280656164

```

# x_cols = list(df_outlier_filter.columns)
# x_cols.remove(outcome)

# train4, test4 = train_test_split(df_outlier_filter)

# predictors = '+'.join(x_cols)
# formula = outcome + '~' + predictors
# model4 = ols(formula=formula, data=train4).fit()
# model4.summary()

X4 = df_outlier_filter.drop(columns=['price'], axis=1)
y4 = df_outlier_filter['price']
X_train4, X_test4, y_train4, y_test4 = train_test_split(

# Refit model with subset features
predictors = sm.add_constant(X_train4)
model_4 = sm.OLS(y_train4, predictors).fit()
model_4.summary()

```

OLS Regression Results

Dep. Variable:	price	R-squared:	0.730			
Model:	OLS	Adj. R-squared:	0.730			
Method:	Least Squares	F-statistic:	4678.			
Date:	Mon, 29 Mar 2021	Prob (F-statistic):	0.00			
Time:	14:04:17	Log-Likelihood:	-2162.6			
No. Observations:	17277	AIC:	4347.			
Df Residuals:	17266	BIC:	4433.			
Df Model:	10					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	14.2244	0.210	67.717	0.000	13.813	14.635
bedrooms	-0.0612	0.009	-6.530	0.000	-0.080	-0.042
bathrooms	0.0602	0.005	12.300	0.000	0.051	0.070
sqft_living	0.0002	4.85e-06	40.562	0.000	0.000	0.000
sqft_lot	0.0609	0.003	19.535	0.000	0.055	0.066
floors	0.0602	0.005	11.676	0.000	0.050	0.070
waterfront	0.5191	0.026	20.178	0.000	0.469	0.570


condition	0.0618	0.004	17.621	0.000	0.055	0.061
grade	0.1804	0.003	58.739	0.000	0.174	0.180
yr_built	-0.0018	0.000	-16.730	0.000	-0.002	-0.001
dist_to_dtw	-0.0187	0.000	-73.247	0.000	-0.019	-0.018
Omnibus:	327.656	Durbin-Watson:	1.984			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	559.364			
Skew:	-0.163	Prob(JB):	3.43e-122			
Kurtosis:	3.819	Cond. No.	2.97e+05			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.97e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
fig = sm.graphics.qqplot(model_4.resid, dist=stats.norm,
```



```
regression4 = LinearRegression()
regression4.fit(X_train4, y_train4)

#use the regression for the train and test data
y_hat_train4 = regression4.predict(X_train4)
y_hat_test4 = regression4.predict(X_test4)

#Root Mean Square Error
train_rmse4 = np.sqrt(mean_squared_error(y_train4, y_hat_train4))
test_rmse4 = np.sqrt(mean_squared_error(y_test4, y_hat_test4))

print(f'Train Root Mean Square Error: {train_rmse4}')
print(f'Test Root Mean Square Error: {test_rmse4}')
```

```
Train Root Mean Square Error: 0.27423613934382146
Test Root Mean Square Error: 0.26984535362250917
```

This is a nice improvement. This is our best model thus far. It passes the normality check from looking at the QQ plot and it is homoscedastic.

Interpreting this model:


R-squared: 73.1% variation in the price can be explained by all of our feature columns.

Durbin-watson: A value preferred between 1-2 implies that the regression results are reliable from the side of homoscedasticity.

The highest coefficients belong to Grade and Waterfront: namely, what grade the home has been given by the King County Housing System. Additionally, having a waterfront view as a part of your home largely impacts the price.

When needed, we can now use this model to give us prediction values for an estimated price, given the values for the features of a home we are trying to sell. Obviously, someone would be unable to renovate their home to suddenly have a waterfront view, but doing something like adding a bathroom (the 3rd highest coefficient) seems to also have a significant impact of the expected price of a home for this model as well.


```
model_4.params[1:].sort_values(ascending=False).plot.bar
```



This is a visualization of our coefficients. To compare, I have taken the absolute value of each in the series, but made sure to indicate negative coefficients in red columns. I have also compared the Root Mean Squared Errors of our 4 models below.

```
print(f'Train Root Mean Square Error 1: {train_rmse}')
print(f'Test Root Mean Square Error 1: {test_rmse}')

print(f'Train Root Mean Square Error 2: {train_rmse2}')
print(f'Test Root Mean Square Error 2: {test_rmse2}')

print(f'Train Root Mean Square Error 3: {train_rmse3}')
print(f'Test Root Mean Square Error 3: {test_rmse3}')

print(f'Train Root Mean Square Error 4: {train_rmse4}')
print(f'Test Root Mean Square Error 4: {test_rmse4}')
```

```
Train Root Mean Square Error 1: 0.2742010910106044
Test Root Mean Square Error 1: 1764.4910958009116
Train Root Mean Square Error 2: 0.27395412776536393
Test Root Mean Square Error 2: 0.274568980060669
Train Root Mean Square Error 3: 0.3012951581957556
Test Root Mean Square Error 3: 0.3076316639605217
Train Root Mean Square Error 4: 0.27423613934382146
Test Root Mean Square Error 4: 0.26984535362250917
```

Just from glancing at this, I believe the best model to be Model 4. Although it may be slightly more overfitted than Model 2, Model 4 has the lowest Root Mean Squared Error on its test data. I'll now fit the model on our data, without a train test split.

```
X_final = df_outlier_filter.drop(columns=['price'], axis
y_final = df_outlier_filter['price']

predictors = sm.add_constant(X_final)
model_final = sm.OLS(y_final, predictors).fit()
model_final.summary()
```

OLS Regression Results

Dep. Variable:	price	R-squared:		0.731		
Model:	OLS	Adj. R-squared:		0.730		
Method:	Least Squares	F-statistic:		5851.		
Date:	Mon, 29 Mar 2021	Prob (F-statistic):		0.00		
Time:	14:21:05	Log-Likelihood:		-2632.6		
No. Observations:	21597	AIC:		5287.		
Df Residuals:	21586	BIC:		5375.		
Df Model:	10					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	14.1879	0.187	75.701	0.000	13.821	14.555
bedrooms	-0.0610	0.008	-7.312	0.000	-0.077	-0.045
bathrooms	0.0631	0.004	14.468	0.000	0.055	0.072
sqft_living	0.0002	4.31e-06	44.623	0.000	0.000	0.000
sqft_lot	0.0629	0.003	22.683	0.000	0.057	0.068
floors	0.0590	0.005	12.825	0.000	0.050	0.068
waterfront	0.4995	0.023	21.736	0.000	0.454	0.545
condition	0.0595	0.003	19.086	0.000	0.053	0.066
grade	0.1812	0.003	66.004	0.000	0.176	0.187
yr_built	-0.0018	9.55e-05	-18.582	0.000	-0.002	-0.002
dist_to_dtn	-0.0189	0.000	-82.640	0.000	-0.019	-0.018
Omnibus:	358.329	Durbin-Watson:		1.991		
Prob(Omnibus):	0.000	Jarque-Bera (JB):		582.147		
Skew:	-0.156	Prob(JB):		3.88e-127		
Kurtosis:	3.741	Cond. No.		2.97e+05		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.97e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
regression_final = LinearRegression()
regression_final.fit(X_final, y_final)

y_hat_final = regression_final.predict(X_final)
rmse_final = np.sqrt(mean_squared_error(y_final, y_hat_f

print(f'Test Root Mean Square Error: {rmse_final}')
```

Test Root Mean Square Error: 0.2733395685228193

Conclusion

I believe the best model is Model 4, where the outliers have been filtered out and none of the features are removed . Although this suffers from multicollinearity, it has an r-squared value of ~0.73, which is the most accurate model in our analysis.

I believe this is acceptable within the context of this scenario. It affects the coefficients and p-values, but it does not influence the predictions, precision of the predictions, and the statistics determining goodness of fit. Our primary goal is to have a model to make predictions for us.

To further improve this, I would use more of the columns included in the original dataset to try to increase my r-squared value and hopefully fix the QQplot issues I was having for all of my models.